MOSES Flight Software
A Guide to the Source

David Keltgen, Roy Smart, Jackson Remington

May 25, 2015

Contents

1 Revision History

Revision	Date	History	Initial
0.1	03-11-2015	Created	DJK
0.2	4-26-2015	Updated	RTS
0.2	4-30-2015	Updated	JRR

2   Introduction

The MOSES instrument is a sounding rocket payload designed to take exposures of the Sun in Extreme Ultraviolet wavelengths. Since the sounding rocket trajectory guarantees only 5 minutes of viable exposure time, the flight software must be able to operate nearly autonomously for the entire flight. Additionally, the flight software must also be able to be controlled from the ground to be able to integrate and test the instrument properly.

MOSES first launched in 2006 using flight software developed by Reginald Mead in C++, running on a Hercules EBX flight computer. Unfortunately this old flight computer started to develop electrical problems and could not be replaced. These events necessitated the selection of a new flight computer and development of new flight software. The new configuration was designed very similar to the original configuration in that an FPGA is utilized to capture the experimental data from the cameras, and that data is transferred back to the flight computer using Direct Memory Access. As a result the flight software is similar to the software developed by Reginald, except that it is written in C.

While the old flight software certainly performed as expected on the 2006 flight, the developers of the updated code made many attempts to fix or implement features that were not implemented in time for the first flight. These include: reducing the latencies between subsequent exposures, a well tested telemetry module, and good integration between ground station software and flight software.

With all that being said, the developer must note that this document is not intended as a Software User's Guide. It is designed to be a resource for the maintainer of the MOSES flight software and to provide a source code level description of the program. While all attempts will be made to make this document as accurate as possible, there will inevitably be errors that creep into this writing. The source code is the main source of information on the operation of the flight software, and an in-depth understanding of the software will only be possible through reading the source.

3   Main Subsystems

1. Flight Computer (FC) stack: The FC stack consists of four boards that work together to capture scientific data and control the instrument.

   (a) Tri-M VDX104+: Main flight computer board. Contains the CPU, SD hard drive, RS-232 ports, and PCI bus. This is the board that executes the flight software.

   (b) CTI FreeForm PCI-104 Virtex 5 FPGA: Connected to the VDX104+ through PCI bus. Main function is to capture science data produced by the ROE. Also implemented on the FPGA is all of the output/input GPIO lines required for operation of power subsystems (outputs) and Timers/Uplinks (inputs).

   (c) Synclink USB Adapter: Connected to the VDX flight computer. Its job is to send science data using RS-422 protocol to the telemetry section at 10 Mbit/s.

   (d) Fourth?

2. MOSES Instrument Electronics: The MOSES instrument relies on several electronic subsystems to achieve it's goal of capturing images of the Sun. These systems operate mostly independently, but rely on the FC stack for control.

   (a) Charged-Coupled Devices (CCDs): MOSES employs a diffraction grating as its primary optical element to provide data in three spectral orders. The instrument has three CCDs that capture the images from each order independently. Multiple coatings on the grating are able to block out (nearly) all wavelengths besides the target lines of 459 and 465 Angstroms. Each exposure creates three different images with uniquely-shifted overlap of the target lines. These spectral orders are known as Minus, Zero, and Plus; together, they can be used to generate a three-dimensional representation of solar activity.

   (b) Read-out Electronics (ROE): The ROE is the interface between the FPGA and the three CCDs that capture scientific data. Upon receiving a command from the flight computer, it begins readout by clocking the data contained in each pixel of the CCDs and streaming it in serial to the SPU (defined below). The FPGA then buffers and formats converted data into separate channels that can be identified by the flight software. The ROE was not designed specifically for MOSES, but for the Hinode instrument which has four CCDs. As a result, the ROE actually provides a fourth data channel: Noise. Noise is labeled the 0 channel while the Minus, Zero and Plus are respectively labeled 1, 2, and 3. (?????Check!!!)

   (c) Serial-to-Parallel Unit (SPU): The SPU converts serial data from the ROE into 16-bit, 32 Mbit/s parallel data stream that will be captured by the FPGA.

(d) Power board: This board manages the power systems on the instrument. The flight computer applies a 'high' value to whichever power subsystems have been requested to be activated and then strobes a latch to turn it on/off.

(e) High-Speed Telemetry Transmitter: Science data is sent from the FC, through the Synclink USB adapter, Premod filter, and finally to a radio transmitter which sends the data to the ground at 10 Mbit/s.

(f) Timers and Uplinks: Timers and Uplinks are single-ended GPIO lines that are used to provide basic control to the instrument. When triggered, they consist of a 5V rising edge.

(g) Housekeeping Link Protocol (HLP): Consists of two separate radio connections, HKUP and HKDOWN. HKUP operates at 1200 baud and sends uplinks from the ground to the instrument. HKDOWN operates at 9600 baud and sends replies from the instrument to the ground. The HLP link is the main way of controlling the instrument during testing and provides the most control over the instrument.

(h) Shutter: The shutter opens and closes (obviously) to allow light to reach the CCDs for each exposure. The this operation is controlled by two GPIO lines connected to the VDX flight computer. The reason these lines are separate from the rest of the GPIO lines implemented through the FPGA is the developers felt that the FPGA may be busy during the time which the shutter is intended to close.

3. Electronic Ground Station Equipment (EGSE): The EGSE is designed to interpret and display data sent back to the ground. It consists of three computer systems that manage the different types of data provided by the instrument. These systems are located in a server tower with the necessary electronic support as well as power supplies to operate the instrument during testing.

(a) EGSE1: The computer labeled EGSE1 is a computer running Microsoft Windows XP that is designed to capture analog housekeeping data produced by the instrument. The software that captures this analog data is written in LabView and provides graphs of temperatures and currents on the instrument over time.

(b) EGSE2: The computer labeled EGSE2 is a computer running (Linux Mint 17???) used for running the EGSE software. This software is a Java program consisting of a Server and Client that work together to exchange HLP packets with the flight computer. EGSE2 is also connected to the flight computer via an ethernet connection while the instrument is on the ground. This allows users to open an SSH session with the flight computer for debugging purposes.

(c) EGSE Laptop: This Ubuntu 14.04 computer runs a program known as receiveTM to capture high-speed telemetry sent by the MOSES instrument. Received data is then piped to an IDL program which displays the images.
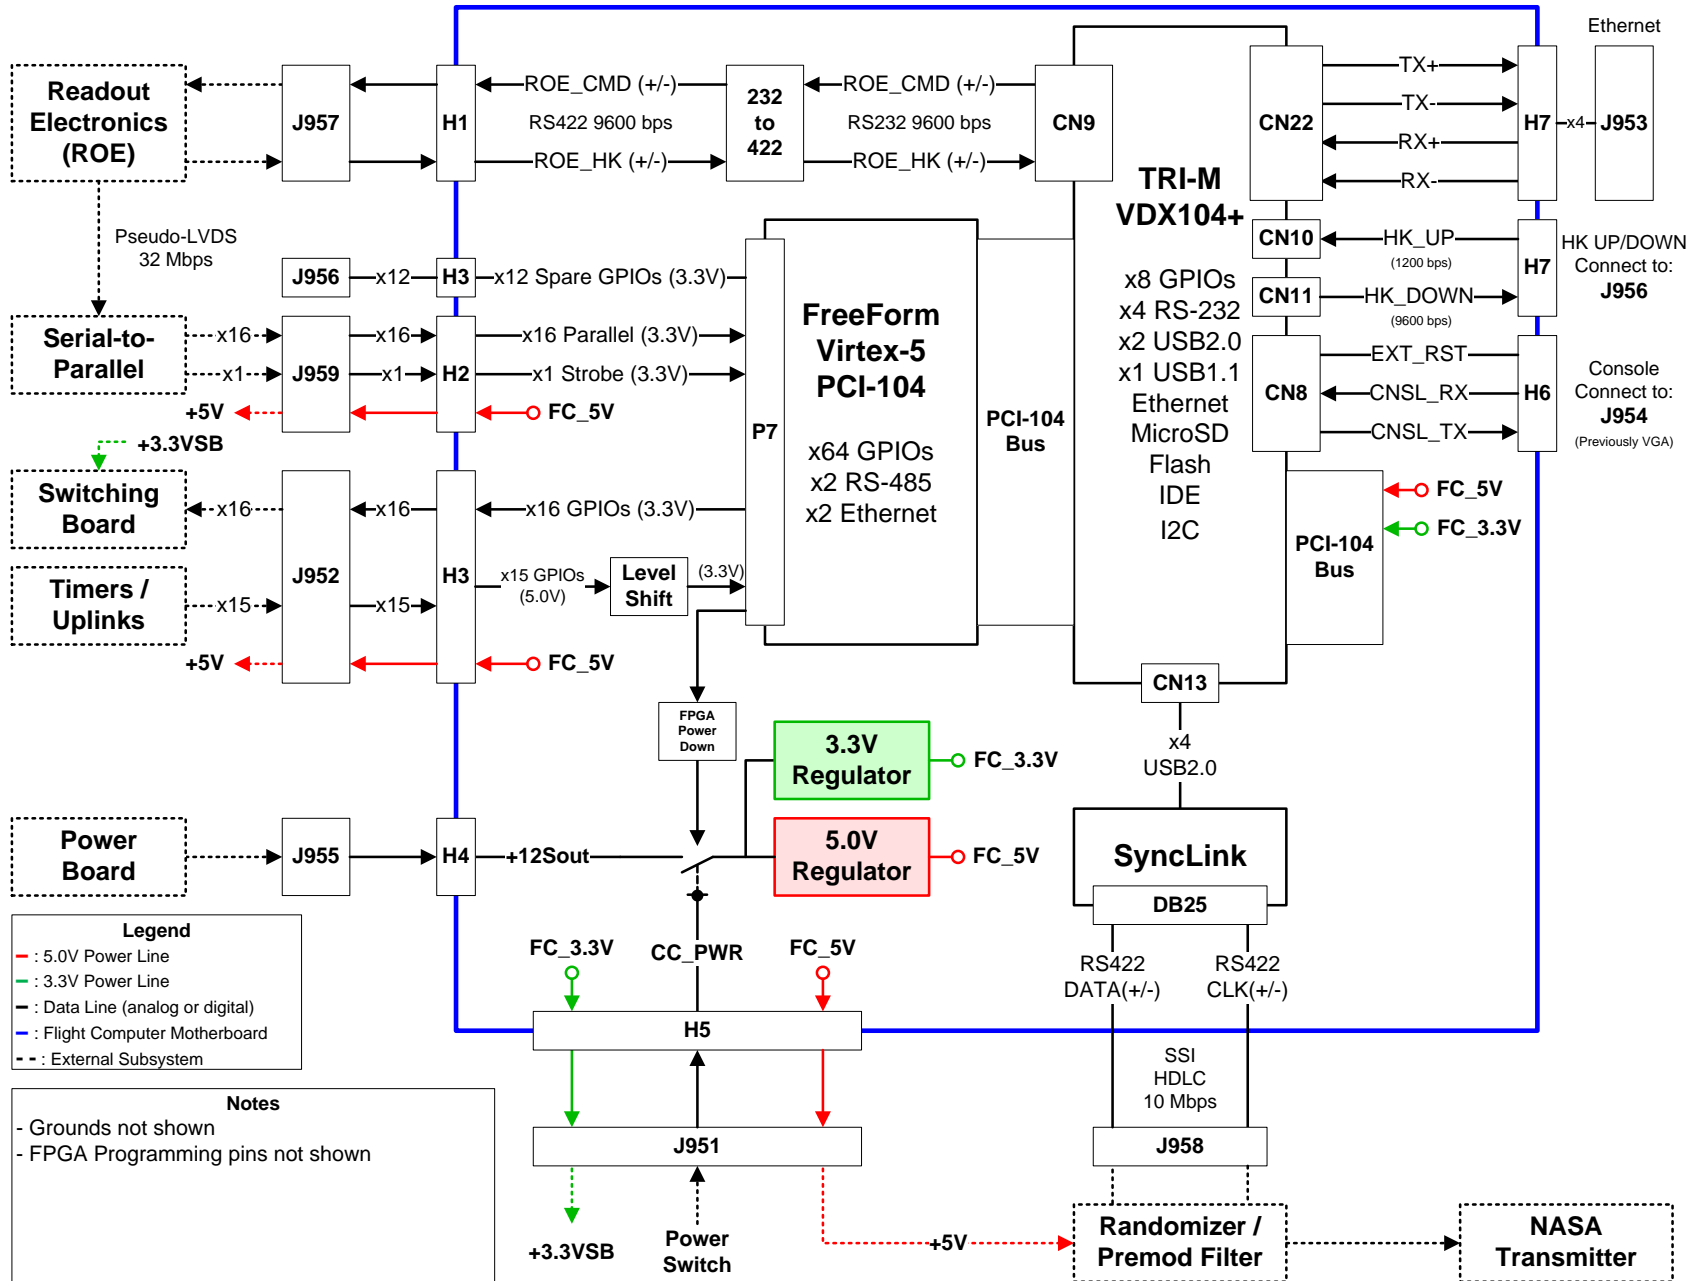
Readout Electronics (ROE)

J957 H1

ROE_CMD (+/-) RS422 9600 bps

ROE_HK (+/-)

232 to 422

ROE_CMD (+/-) RS232 9600 bps

ROE_HK (+/-)

CN9

CN22

TX+
TX-
RX+
RX-

H7 x4 J953

Ethernet

TRI-M VDX104+

Pseudo-LVDS 32 Mbps

J956 x12 H3 x12 Spare GPIOs (3.3V)

Serial-to-Parallel

x16 J959 x16 H2 x16 Parallel (3.3V)
x1 x1 x1 Strobe (3.3V)

+5V FC_5V

+3.3VSB

Switching Board x16 J952 x16 H3 x16 GPIOs (3.3V)

Timers / Uplinks x15 x15 x15 GPIOs (5.0V) Level Shift (3.3V)

+5V FC_5V

FreeForm Virtex-5 PCI-104

P7

x64 GPIOs
x2 RS-485
x2 Ethernet

PCI-104 Bus

CN10 HK_UP (1200 bps) H7

CN11 HK_DOWN (9600 bps)

HK UP/DOWN Connect to: J956

CN8 EXT_RST
CNSL_RX
CNSL_TX H6

Console Connect to: J954 (Previously VGA)

x8 GPIOs
x4 RS-232
x2 USB2.0
x1 USB1.1
Ethernet
MicroSD
Flash
IDE
I2C

FC_5V
FC_3.3V

PCI-104 Bus

CN13

FPGA Power Down

3.3V Regulator FC_3.3V

5.0V Regulator FC_5V

x4 USB2.0

SyncLink

DB25

Power Board J955 H4 +12Sout

RS422 DATA(+/-) RS422 CLK(+/-)

Legend
— : 5.0V Power Line
— : 3.3V Power Line
-- : Data Line (analog or digital)
— : Flight Computer Motherboard
-- : External Subsystem

FC_3.3V CC_PWR FC_5V

H5

SSI HDLC 10 Mbps

Notes
- Grounds not shown
- FPGA Programming pins not shown

J951

J958

+3.3VSB Power Switch +5V

Randomizer / Premod Filter

NASA Transmitter

Figure 1: Block diagram summarizing the hardware architecture for interfacing with the MOSES instrument.

4 Requirements for Operating the MOSES Instrument

1. Operate power systems on the instrument. The power systems include:

   (a) Shutter Driver: Must be activated to control the shutter.

   (b) ROE: Must be activated to readout and flush the CCDs for exposures.

   (c) Premod Filter: This system must be activated for high-speed telemetry to be used. It uses a hardware random number generator to randomize the telemetry data. This is important as the telemetry stream must have equal numbers of ones and zeros to be properly reconstructed on the ground.

   (d) Temperature Control Systems: These systems provide cooling to the instrument while under vaccum. The CCDs are cooled using this system to minimize the noise inherent in the CCDs. The flight computer must also be cooled under vacuum to prevent overheating, as it is normally air-cooled. It must be noted that this system is NOT to be turned on while the instrument is at atmospheric pressure.

   (e) 5V Regulator: This must be on to provide the premod filter with a -5V rail.

   (f) 12V Regulator: Among other things, this subsytem activates the analog telemetry monitoring.

   (g) H-α Camera: This camera is used in flight as a targeting system, to ensure the attitude of the payload is appropriate. This is an analog system, so the data produced by the H-α camera is not moderated by the flight computer, it is only activated through the flight computer.

   The flight software is directed to turn on/off power systems through HLP packets sent by the EGSE.

2. Control the read-out electronics (ROE). This is accomplished using an RS-422 serial connection between the FC and the ROE. Common task include: commanding exposures, flushing the CCDs, and requesting housekeeping data such as voltages and currents.

3. Open and close the Shutter. The Shutter depends on two GPIO lines (open and close) and each must be pulsed for 200ms to execute their respective operations.

4. Receive and save science data. Through the SPU, the ROE sends science data over a 32 Mbit/sec parallel connection. The FSW should have low enough latency so as to not miss any science data. The FSW should save each image as a 16 MB file with the extension .roe.

5. Create an index of the images that were captured. This index is critical to the IDL software used to analyze MOSES images. This index should contain the name of the file, the number channels in the file

6. Send science data over telemetry. NASA has provided a 10 Mbit/sec serial line that connects with a radio to the ground. During the course of the mission.

While there is not enough time to send all of the data, the MFSW should send as much science data back to earth as possible to mitigate data loss from the harsh re-entry environment.

7. Respond to Timers and Uplinks. Timers are single lines provided by NASA that instruct the instrument when to execute the different parts of the experiment. Timers include:

   (a) Dark Exposure Start: These exposures don't open the shutter while taking an exposure. They are used to provide a baseline for the CCDs for data analysis post-flight.

   (b) Data Start: This command carries out the data sequence outlined in a sequence file (to be explained later).

   (c) Data Stop: Stops the current sequence whether it be a dark sequence or a data sequence. The sequence stops only after the current exposure has completed. Note: This command only stops the current data sequence. If there is more than one data sequence enqueued to the Science Timeline, the FSW will just start the next sequence following Data Stop.

   (d) Sleep: Instructs the flight software to stop and to shut down the computer. This is important as we don't want the flight computer to be damaged during reentry into the atmosphere.

   Dark exposures are those which don't open the shutter while the CCDs are being exposed during the Science Timeline. Uplinks are similar to Timers in that they have the same functionality, except they are tied to a physical button on the ground. Using this interface, a user could control the experiment from just the push-button on the EGSE tower.

8. Respond to HLP packets. These packets are sent by the EGSE software on the ground, and provide all the functionality of Uplinks and Timers while providing additional commands to control the instrument. Two RS-232 connections are provided by NASA to facilitate this interface: HKUP and HKDOWN. Possible types of HLP packets include:

   (a) Uplink: Provides the same functionality as Uplinks initiated through the GPIO interface (e.g. Data Start, Data Stop, etc.)

   (b) Shell: The FSW opens a bash shell as a child process. Shell packets can execute bash commands on the FC via this interface.

   (c) Power: Queries, activates or deactivates power subsystems.

   (d) Housekeeping Requests: Requests housekeeping data from the ROE.

   (e) Mission Data Acquisition Requests: Allows direct control over exposure parameters and can also control the ROE.

   The HLP interface is the main interface used for debugging and testing the instrument. This interface is also important during flight as it sends packets that inform the users of the current state of the FSW.

5 Architecture Description

The MOSES instrument needs to be able to respond to IO on several different interfaces. This is problematic for a sequential programming architecture, as input could be lost while the computer is executing another part of the program or output could be executed late while the software is waiting for something else. This issue is alleviated by using a threaded software architecture, which can execute separate subprograms concurrently. Linux provides excellent support for threaded programming, known as POSIX threads (pthreads). Pthreads allows the flight software to execute the Science Timeline, while still being available to respond to commands from the ground or write data to the hard-disk drive.

The challenge with this threaded architecture is one of thread synchronization. Each thread operates independently, and steps must be taken to ensure that the program executes in the proper order and no memory is accessed simultaneously by two or more threads. In the MSFW thread synchronization is accomplished through so-called Locking Queues. These objects are implemented anywhere in the porgram where the producer-consumer problem is present, where one thread is producing data while the other is waiting to do operations on that data. Locking Queues are identical to a normal queue data structure, except that they take advantage of mutual-exclusion (mutex) locks to "lock" the queue until the accessing thread has completed its operation on the queue. The Locking Queue also uses a conditional variable to get the attention of the consumer waiting for input. Another thread synchronization technique utilized in the MFSW is signals. Signals are objects provided by the Linux OS that allow separate threads or processes to get each others attention through a binary flag. In the MFSW, the SIGINT signal is used to instruct the main process to shut down the flight software. Another signal, SIGUSR1 is used as a backdoor to command exposures during FSW debugging. This allows rudimentary control over the FSW without the need for the EGSE software. Finally, the last thread synchronization technique is known as a semaphore. Only one semaphore is used in the MFSW, and it is used to synchronize operations between the FPGA Server thread and the Science Timeline thread. A Locking Queue would have been used here, but as of this writing there seems to be a bug in the `pthread_cond_timedwait()` function that prevents it from operating properly.

The software is broken up into threads based off of input/output requirements. For the most part, each thread represents one IO interface that can only be controlled by the associated thread. The exception to this rule is the ROE CMD/HK interface, which is accessed by both the Science Timeline and the HLP Control threads. Upon program start, the first thread to be executed is the Main thread. Its purpose is to start all the other threads and wait for a signal to shut down the FSW. The most important thread in the software is the Science Timeline thread. This thread controls the timing of the experiment, while relying on other threads to communicate with the appropriate interfaces. Another important thread is the FPGA Server thread. This thread mediates communications between the FSW and the FPGA, and is directly responsible for capturing science data and for notifying the software of GPIO input. All of the other threads are usually responsible for separate IO interfaces and will be explained in-depth below.
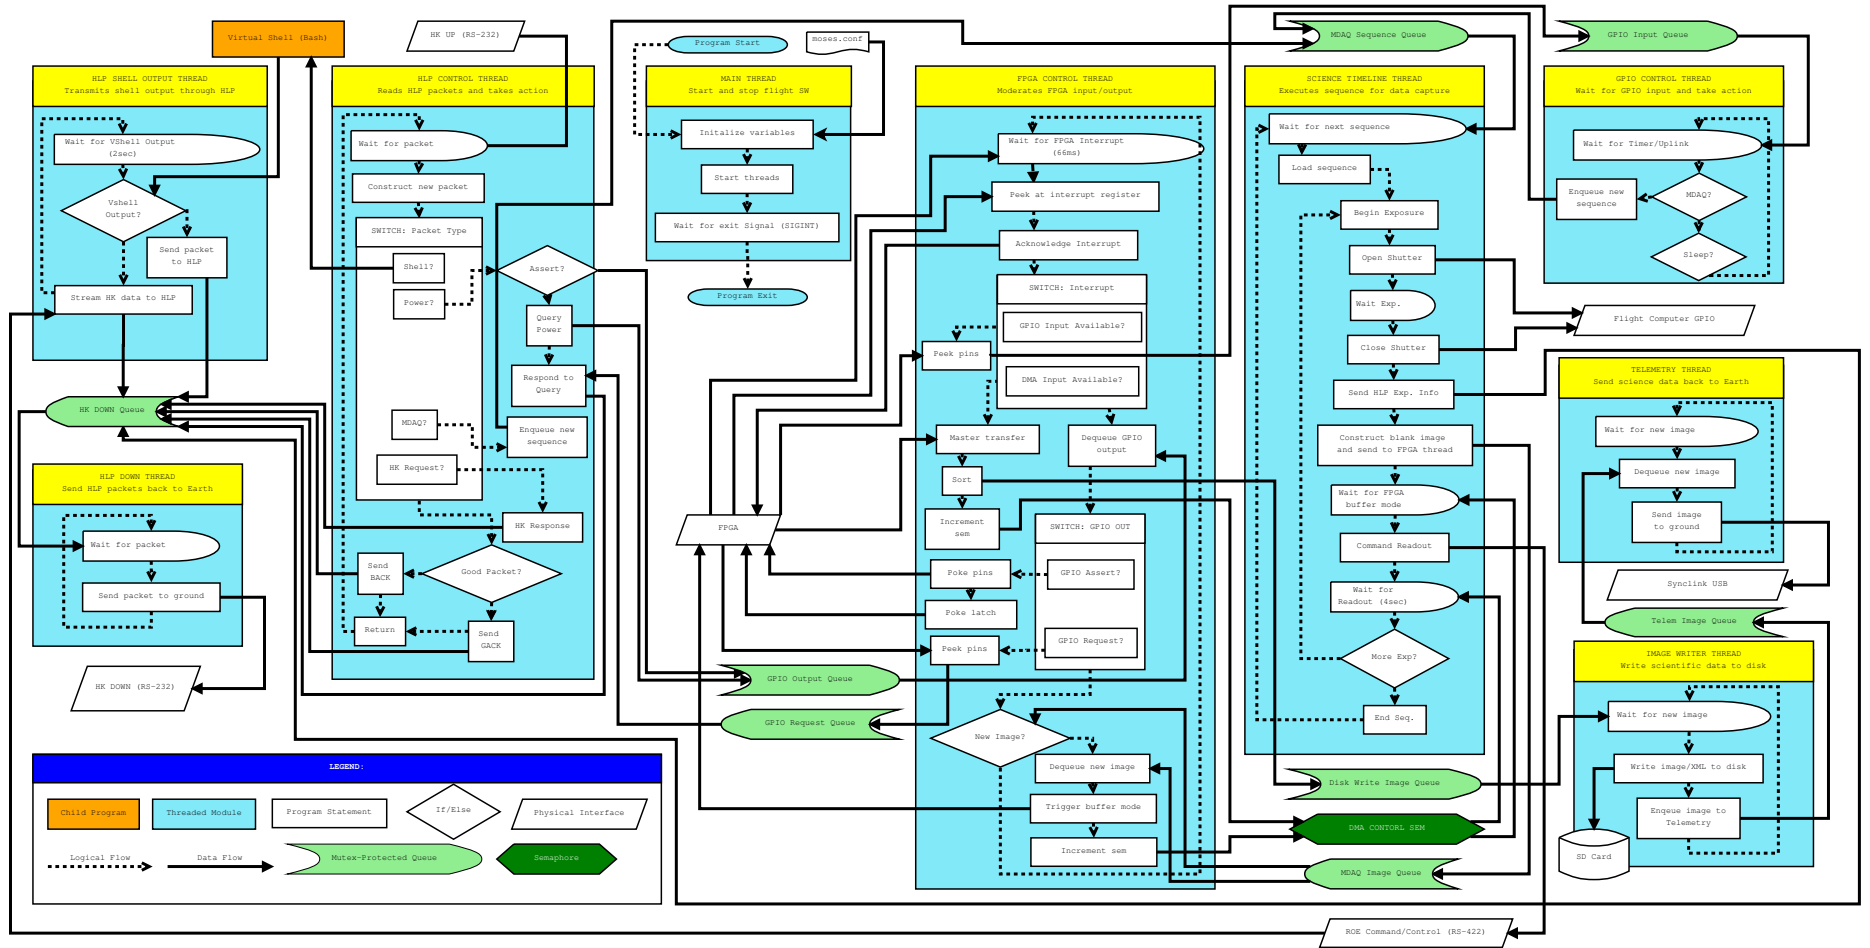
Figure 2: Block diagram summarizing the flight software architecture.

6   Main Thread

Program initialization begins with the configuration file `moses.conf`. Each setting in the file specifies an individual thread or method of communication and is assigned a one (ON) or zero (OFF) for toggling during the debug process. After reading the configuration file, the Main thread starts the `bash` subprocess and notes its Program ID (PID). This will be used to kill the `bash` subprocess upon program exit. Then the Main thread proceeds to start each thread specified as active by `moses.conf`. The order of threads in `moses.conf` designates their load order; i.e. items listed higher will have higher execution priority. During flight, every option shall be set to ON.

After each thread has successfully started, the Main thread is paused and simply waits for program exit. To exit the program any one of the threads may send a `SIGINT` Linux signal to the main thread to inform it of program shutdown. The main thread is waiting for this signal, and once it is raised, it will begin the process of cleaning up threads and exiting the program.

7 Science Timeline Thread

As the name implies, the Science Timeline thread controls the overall experiment schedule. It starts off by ensuring that the ROE is active, otherwise the taking exposures would be pointless. Like all of the threads, Science Timeline is initialized at startup, but waits on a signal ('SIGUSR1') in order to continue. Once it has been received, the thread will set the current sequence from the already-initialized sequence map and begin taking exposures. The 'takeExposure' function starts off by flushing the ROE Charge-Coupled Devices (CCDs) 5 times in order to remove any accumulated charge. It then opens the shutter and then uses a sleep function to wait for the designated exposure time (timing experiments to determine accurate sleep times are still to be done). If the exposure is a data sequence, then additional time will be added to the exposure length as represented by the variable PULSE in 'sciencetimeline.h.' This additional time is used to account for the difference in the time it takes to open and close the shutter. Once the exposure is taken, meta data of the image is saved to the data structure. After that, the thread commands the ROE to read out its exposure data. This process will take four seconds, and this thread can sleep during that time (to free other processes). Science Timeline's last task to complete is to enqueue the image to the image-writing thread to independently start writing data to the SD card while the Science Timeline thread moves to the next exposure. This process is repeated for each exposure in the specified sequence.

7.1 Read Out Electronics (ROE)

The ROE is capable of being in several different modes. These modes are:

Default Mode: This is the mode that the ROE is initially in on startup. In default mode, only one command can be sent to it, which is the exitDefault() command. While in default mode, all the ROE does is simply readout an exposure every twelve seconds.

Command Mode: This is the mode that the ROE is in after the exitDefault() command has been sent. The ROE will be in command mode for the majority of the flight. In this mode, commands can be sent such as setting the ROE to a new mode, reading out the CCDs or requesting Housekeeping values.

Selftest Mode: In selftest mode, the ROE will not read out the CCDs at all, but will instead read out a predefined sequenceof vertical bars. This can be useful for testing the flight computers data acquisition abilities. Once the ROE is in selftest mode, it must be reset via the reset() method in order to exit selftest mode. One would also need to exit default mode again after the reset.

Stims Mode: In stims mode, a predefined pattern is generated in much the same way as the selftest pattern. The difference is that in stims mode the pattern is fed through the CCD readout circuits and thus any extraneous noise and other anomalies show up in the pattern. Also, unlike selftest mode, stims mode can be exited by using the stimOff() method.

Nominal Operation

The ReadOutElectronics object is instructed to exit default mode, reset, and then exit default mode again. The cause for this seeming redundancy is that commanding the ROE to exit default mode while in command mode has no effect. However, commanding the ROE to reset while it is in default mode causes a software error. Therefore, we exit

default mode once to make sure that the ROE is commandable, tell it to reset so that we know what configuration it is in, and then finally exit default mode again to make it commandable in a known configuration. The default port for the ROE is /dev/ttyS1 (COM2). The main job of the ROE is to read out the data on the CCD cameras. The two main functions required to acheive this are the flush() and readout() functions. Flush is used to clear any accumulated data on the cameras and readOut is used to get the data from the cameras. The biggest difference between the two, is that flushing the ROE is faster and doesnt send any data down the data link. It is recommended that the ROE be flushed five times in succession before every exposure. Readout does just what one would expect, it reads out the cameras. It is suggested that the ROE be allowed four seconds for reading out data between exposures. The only notable point of confusion for this method is that it also requires a block id, which is basically just a byte that defines how the ROE needs to be read out. There are unique ids for reading out the ROE normally, reading out while in selftest mode, reading out while in stims mode, and there are even some undocumented ids which were used for testing during the construction of the ROE. All of these block ids are documented in the roe.h file.

As a last point of interest, there are also methods for getting and setting the analogue electronics parameters inside the ROE. These methods normally go unused and are usually only useful during testing. These parameters consist of 8 bytes which control the ROEs behavior. In fact, it is by setting these values that the ROE is placed into selftest and stims mode. In those methods these parameters are written automatically. It is suggested that these parameters only be changed if an experienced ROE operator /technician knows what they are doing.

In regards to the Housekeeping data of the ROE, the values that are returned are the raw hex values. These raw hex values are then passed on from the flight software to the ground station. The ground station then displays these raw values, unconverted.

Mutex locks have been implemented in the ROE functions since multiple threads could potentially be trying to access it at the same time. An example of this could be that the Science timeline thread could be telling the ROE to flush the CCDs, while a packet sent from the GSE could be requesting HK values. Although the likelihood of this clashing of threads causing errors in small, it is important to make sure that the signals being sent to the ROE are exactly what we intended them to be.

8   Image Writer Thread

Through time testing, it was determined that previous revisions took approximately one second to write each ROE image to disc. Since this mission is very time-critical, such a delay was deemed too long to fit serially in the Science Timeline thread. Moving it to another results in a large percentage of flight time saved, allowing more exposures to be captured per flight. A ROE image structure is used to store all the image data and meta data for each successful image. Here is a brief overview therein:

Identifying values for the origin, instrument, observer, and object will typically not change on each flight. However, if this software were to be used under different circumstances, these values could be changed to reflect that. 'Bitpix,' the bits per pixel value, along with the images' width and height, are set...(???) 'Channels' consists of a single character value that can be ANDed and ORed with defined values to enable or disable each channel as required. For ideal exposures, channels 1, 2, and 3 will be enabled while channel 0 (noise) is disabled. Other values unique to each image include the filename, its path, the date and time its exposure started, its duration, the size in pixels of each channel, as well as the total number of exposures in its sequence.

The image writer thread starts by assigning data to the image structure inside the 'write_data' function. Once all of these values are assigned, it goes to write the file. This is subsequently broken down into two tasks: writing the actual image data to disk, and writing the record file (imageindex.xml) to disc. The xml file contains all of the metadata.

Rather than be discarded after saving, the files are kept in memory (aside from the noise channel, which is only saved locally). A pointer to it is placed into the telemetry queue for transmission to the ground. This way, many images may be retrieved before touchdown and without requiring survival of the payload. An updated version of the xml file is also sent to the queue between each image.

9 GPIO Control Thread

The GPIO Control thread is the primary means of controlling the software during flight. After launch, Timers will be raised by the NASA telemetry section. These Timers are set to inform the experiment when to begin Dark sequences, Data sequences, when to stop taking Data sequences, and when to shut the computer down. The Timers are not connected directly to the GPIO Control thread, rather they are mediated by the FPGA Server thread. When the FPGA Server thread receives a Timer, it is enqueued onto the GPIO input queue where it can be interpreted by the GPIO Control thread.

10   FPGA Server

The FPGA Server thread mediates communications with the FPGA and the rest of the flight software. The FPGA is critically important to the goals of the MOSES instrument, as it is the electronics that receives the experimental data from the cameras, as well as managing the GPIO input/output lines which control the instrument and the execution of the flight software. Since there are many different threads that depend on the FPGA for their execution, we implemented the FPGA Server thread to avoid thread contention issues between different parts of the program.

11 Housekeeping Link Protocol (HLP)

11.1 HLP Control Thread

The HLP Control thread listens for incoming HLP packets and then takes the appropriate action based on the content of the packet.

11.2 HLP Down Thread

The HLP Down thread is attached to the HLP Down queue. As the different threads in the flight software produce HLP packets that need to be sent to ground, they enqueue it onto the HLP Down queue. The HLP Down thread dequeues these packets one-by-one and sends it to the ground via a RS-232 serial port.

11.3 Virtual Shell

11.4 HLP Shell Output Thread

12   High-Speed Telemetry

One of the mission requirements for MOSES II is to ensure the retrieval of meaningful data without relying on the payload's survival. The concern arose when MOSES I landed in an active, unexploded ordnance range (luckily, its detonation was not required and the payload was recovered). Such a backup system would even allow for mission success in the event of in-flight failures. To satisfy the requirement, Synclink USB adapters were implemented to manage the high-speed transfer rates of the provided radio transmitter. One is used on each end (flight → ground) to mediate the science data's transfer. Custom programs were implemented in the C language to control the Synclinks; `sendTM`, which is given its own thread in the flight software, and `receive_TM`, a standalone program run by the ground station laptop.

The Telemetry thread is generally initialized with the lowest priority because it has little interference and becomes active last. Its function is to detect objects in the telemetry queue written by the Image Writer thread. When the queue updates and the queue's mutex lock is released, the Telemetry thread accepts the pointer to the next image struct and begins streaming its data to the on-board Synclink. The modules are separated by the Premod filter, which converts the signal to analog for transfer via the provided RF transmitter (as outlined in section 2e).

Meanwhile on the ground, the Ground Station laptop shall be running the receive_TM program, which passively monitors the second Synclink's input. When detecting valid packets, the program displays the size of each, processes the incoming signal, and saves the data to disk. Each full transmission (including all three of the valid channels) takes about 10.1 seconds; each index (.xml) file takes approximately 0.5 seconds, but this value obviously increases with the number of entries recorded within (one for each of the completed exposures). Upon success, the data is saved locally before being piped to an external IDL script, which then reconstructs the images and displays them on screen.

13   Configuring the VDX

The BIOS on the VDX computer must be configured correctly before using the flight software. These steps are only required once each time the CMOS battery is disconnected.

Entering the BIOS on the VDX depends on the board's previous configuration.

- If the VDX board is new, the BIOS will be accessible on COM1. Additionally pin 10 of and COM port must be pulled to ground to enable BIOS access. More information is available in the VDX user manual.

- If the VDX board has been configured before, the BIOS will likely be accessible through the serial console implemented on COM2 it is convenient to have these interfaces on the same serial port.

14   Configuring Serial Ports

The BIOS display can be rather hard to navigate; side-arrow keys switch the active tab and vertical-arrow keys scroll through the menu options of each tab. Unfortunately, through serial port, the cursor will not highlight which option is selected, so the user must keep track.

Open the BIOS by booting the VDX and pressing the delete key. If successful, the terminal window will be taken over by a psuedo-interface surrounded in asterisks. The section tabs are listed along the top border: Main, Advanced, PCIPnP, Boot, Security, Chipset, and Exit. The second line tells you what tab is currently selected. Navigate to the Chipset tab and select SouthBridge Configuration by scrolling to the second option. The menu should appear as in Figure 3. Now scroll to the second option: Serial/Parallel Port Configuration. The new menu should display serial port settings as shown in Figure 4. Ensure the BIOS settings match those shown in the figure.
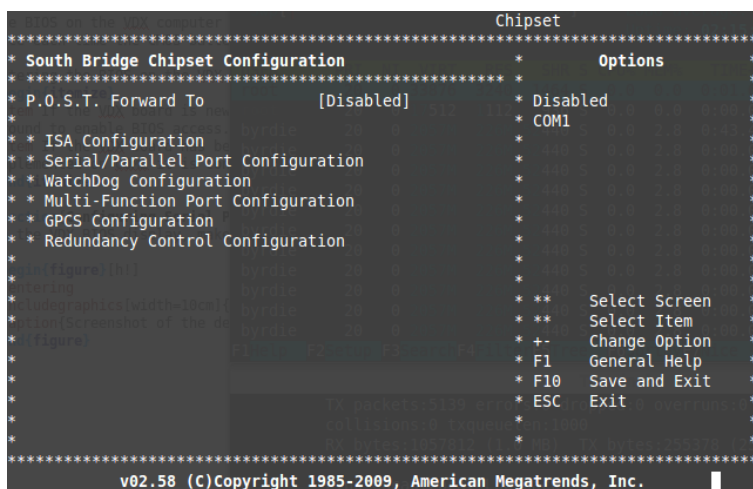

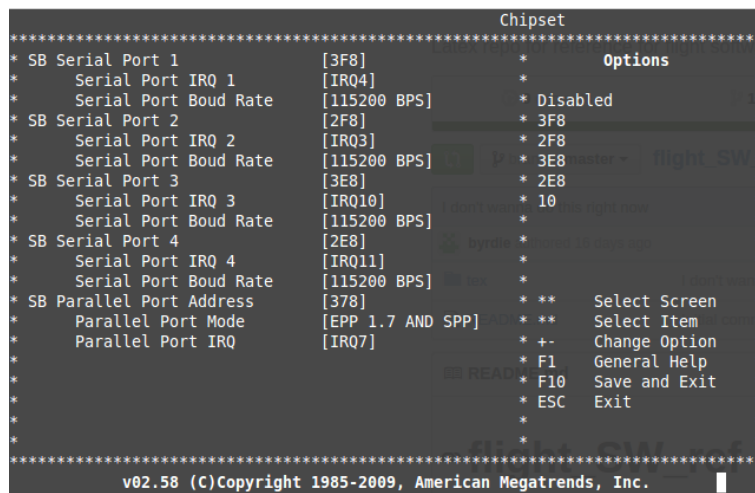Figure 3: Screenshot of the SouthBridge Configuration menu.

```
                                        Chipset
**************************************************************************
* SB Serial Port 1              [3F8]               *          Options          *
*      Serial Port IRQ 1        [IRQ4]              *                            *
*      Serial Port Boud Rate    [115200 BPS]        * Disabled                   *
* SB Serial Port 2              [2F8]               * 3F8                        *
*      Serial Port IRQ 2        [IRQ3]              * 2F8                        *
*      Serial Port Boud Rate    [115200 BPS]        * 3E8                        *
* SB Serial Port 3              [3E8]               * 2E8                        *
*      Serial Port IRQ 3        [IRQ10]             * 10                         *
*      Serial Port Boud Rate    [115200 BPS]        *                            *
* SB Serial Port 4              [2E8]               *                            *
*      Serial Port IRQ 4        [IRQ11]             *                            *
*      Serial Port Boud Rate    [115200 BPS]        *                            *
* SB Parallel Port Address      [378]               * **     Select Screen       *
*      Parallel Port Mode       [EPP 1.7 AND SPP]   * **     Select Item         *
*      Parallel Port IRQ        [IRQ7]              * +-     Change Option       *
*                                                   * F1     General Help        *
*                                                   * F10    Save and Exit       *
*                                                   * ESC    Exit                *
*                                                   *                            *
*                                                   *                            *
**************************************************************************
        v02.58 (C)Copyright 1985-2009, American Megatrends, Inc.
```

Figure 4: Screenshot of the desired VDX BIOS serial port settings.

15 Dictionary of Abbreviations

1. HLP - Housekeeping link protocol

2. ROE - Read-out electronics

3. FC - Flight computer stack (VDX + FPGA)

4. MFSW - MOSES flight software

5. VDX - VDX104+ (Flight computer)

6. GSE - Ground station equipment